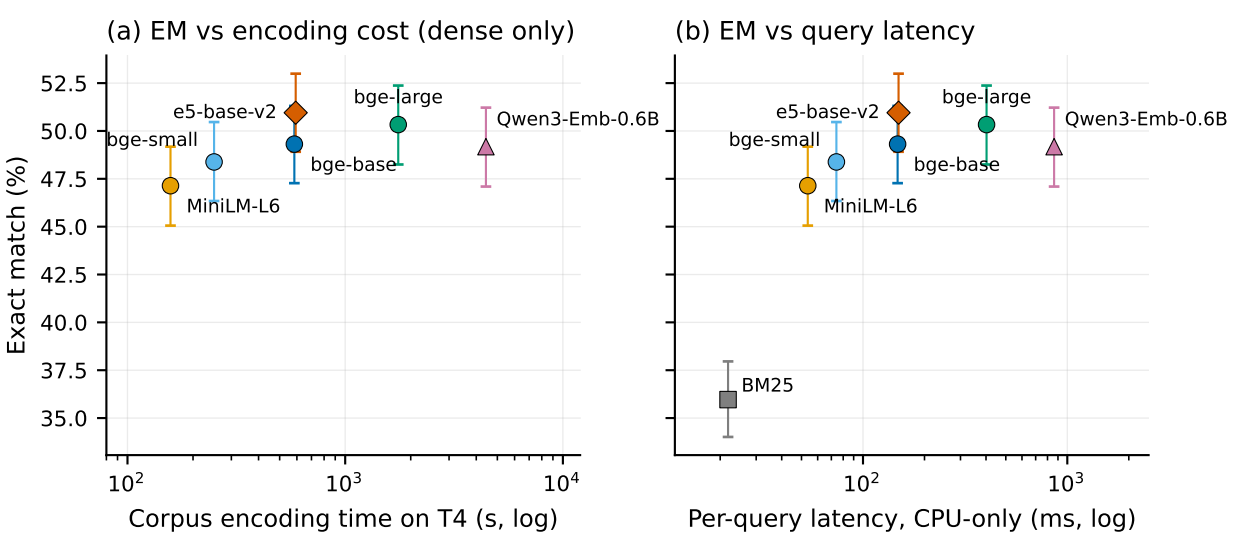
(a) EM vs encoding cost (dense only)
Exact match (%)
52.5
50.0
47.5
45.0
42.5
40.0
37.5
35.0
bge-small
e5-base-v2
bge-large
Qwen3-Emb-0.6B
bge-base
MiniLM-L6
10^2
10^3
10^4
Corpus encoding time on T4 (s, log)
(b) EM vs query latency
bge-small
e5-base-v2
bge-large
Qwen3-Emb-0.6B
bge-base
MiniLM-L6
BM25
10^2
10^3
Per-query latency, CPU-only (ms, log)